

Computer Architecture and system programming

Kristoffer Klokke

2022

Contents

1	2019	5
1.1	Question	5
1.1.1	Convert A and B to binary	5
1.1.2	Interpret A and B as signed integers in two's complement representation. What are their decimal values? . .	5
1.1.3	With the same interpretation of A and B as above, consider their potential multiplication using Booth's algorithm.	5
1.1.4	Still interpreting A and B as signed integers in two's complement representation, perform the subtraction $B - A$ in binary. Indicate the result and argue whether overflow occurred or not.	5
1.1.5	What is the decimal value of C and D ?	7
1.2	Question	7
1.2.1	Write the truth table for F	7
1.2.2	Based on the Karnaugh map, write a minimal Boolean expression for F	7
1.3	Question	8
1.3.1	Describe the address format for using the cache	9
1.3.2	Imagine 6 alternate realities where the next memory request is to read or write to each of the following addresses (written here in hexadecimal). For each scenario, describe whether there is a hit or miss in the cache.	9
1.3.3	The cache has multiple lines per set, so the depiction above does not show all fields. Assume the replacement algorithm used in the cache is LRU, which maintains an additional bit per cache line. Compared to the actual data stored in the cache, how many percent additional memory is used in the cache for all required meta-data?	9
1.3.4	We would like to use this cache design for the cache in each processor in an symmetric multi-processing system. However, the caches are to be kept coherent using the MESI protocol. Which changes to the cache organisation are needed, if any?	10
1.4	Question	10

1.4.1	The program would like to send the following 8 bits of data, written from left to right as d_1, d_2, \dots, d_8 : 1111 0001. Compute the corresponding 12-bit Hamming code, and indicate the check bits.	10
1.4.2	The program is receives the following 12 bits in response, written from left to right as b_1, b_2, \dots, b_{12} : 1001 1100 0001. Assuming that at most 1 single-bit error has occurred, compute the 8-bit message.	10
2	2018	11
2.1	Question	11
2.1.1	Interpret S and T as unsigned integers. What are their decimal values?	11
2.1.2	Interpret S and T as signed integers in two's complement representation. What are their decimal values?	11
2.1.3	Interpret S and T as signed integers in two's complement representation. Multiply S and T using Booth's algorithm. Show	11
2.1.4	What is the range of integers are representable in this format?	12
2.1.5	Which bit pattern represents the lowest and highest representable values?	12
2.1.6	Give the decimal value of A and B in the normalised form $S \cdot 2^E$ where $1 \leq S \leq 2$	13
2.1.7	Perform the addition of A and B in binary arithmetic	13
2.2	Question	13
2.2.1	As a first solution, write a Boolean expression for F in sum-of-products form.	14
2.2.2	Draw the Karnaugh map for F	15
2.2.3	Based on the Karnaugh map, write a minimal Boolean expression for F	15
2.3	Question	15
2.3.1	How many lines will the cache have?	15
2.3.2	Describe the address format for the cache?	15
2.3.3	Besides the actual data, for each line we must also store a tag and a valid-bit. Compared to the memory used for data, how many percent additional memory do we need for the required meta-data?	16
2.3.4	How many sets will the cache have?	16
2.3.5	Describe the address format for the cache?	16

2.3.6	Besides the actual data, for each line we must also store a tag and a valid-bit. Compared to the memory used for data, how many percent additional memory do we need for the required meta-data?	16
2.3.7	The direct-mapped cache has an access time for each memory request of 4ns and a miss rate of 5%. After a miss the request goes to main memory, which adds 103ns to the access time (including the subsequent cache update). What is the average access time?	16
2.3.8	Due to its increased complexity the set-associative cache has a slightly higher access time of 6ns, though the miss rate drops to just 3%. As the cache update in case of a miss is also slightly slower, the added access time is now 104ns. What is the average access time for this setup?	17
2.4	Question	17
2.4.1	For the following sequence of events, show the resulting MESI state of each block in each processor.	17
2.5	Question	18
2.5.1	To what values do the following expressions evaluate	18
2.5.2	What is a void pointer and what is such a pointer used for?	18
2.5.3	Declare a function pointer “ptr” to a function requiring two integers as arguments returning a pointer to a long variable.	19
2.5.4	Give the stack-frame relevant for the function relative to the function’s stack base pointer (%ebp). Identify all variables and parameters of the C-Function. Use the following table as template (produce a table similar to this one):	19
2.5.5	Identify the lines in the assembly code where the abort condition $i < d$ is checked and where does the code continue in case the condition does not hold and where in case it does not.	20
2.5.6	Identify in which lines in the assembler code the address of $a[i]$ is calculated and in which register is the value of $a[i]$ temporarily stored.	20
2.5.7	Consider line 16 in the assembly code	21
2.5.8	Through which register is the result returned to the caller?	21

1 2019

1.1 Question

While decompiling a program you encounter the following hard-coded bytes, shown in hex-adecimal:

$A = 0x6d$

$B = 0xb3$

1.1.1 Convert A and B to binary

$A = 01101101$

$B = 10110011$

For the conversion each value of the hex number can be directly converted to its 4 bit equivalent.

1.1.2 Interpret A and B as signed integers in two's complement representation. What are their decimal values?

$A = 64 + 32 + 8 + 4 + 1 = 109$ $B = -128 + 32 + 16 + 2 + 1 = -77$ Using the two's complement the binary value can be converted to $A = 109$ and $B = -77$

1.1.3 With the same interpretation of A and B as above, consider their potential multiplication using Booth's algorithm.

Result = 01100100

- How many rounds would be performed during the algorithm? - 8 rounds
- How many additions/subtractions would be needed? - 5 operations for both ways of multiplying

1.1.4 Still interpreting A and B as signed integers in two's complement representation, perform the subtraction $B - A$ in binary. Indicate the result and argue whether overflow occurred or not.

$\overline{A} = 10010010$

$B + (\overline{A} + 1)$

1011 0011

1001 0011

10100 0110

First the value of A is negated and one is added to get the actual negated

B as multiplier					
qn	q[n+1]	BR	AC	QR	sc
		initial	0000 0000	1011 0011	8
1	0	A = A - BR	01001010		
		rightShift	0010 0101	0101 1001	7
1	0	A = A + BR	10010010		
		rightShift	1100 1001	0010 1100	6
0	0	rightShift	1110 0100	1001 0110	5
0	0	rightShift	1111 0010	0100 1011	4
1	0	A = A - BR	00111100		
		rightShift	0001 1110	0010 0101	3
1	0	A = A + BR	10001011		
		rightShift	1100 0101	1001 0010	2
0	0	rightShift	1110 0010	1100 1001	1
1	0	A = A - BR	00101100		
		rightShift	0001 0110	0110 0100	0

A as multiplier					
qn	q[n+1]	BR	AC	QR	sc
		initial	0000 0000	0110 1101	8
1	0	A = A - BR	00110011		
		rightShift	0001 1001	1011 0110	7
0	0	rightShift	0000 1100	1101 1011	6
1	0	A = A + BR	10111111		
		rightShift	1101 1111	1110 1101	5
1	0	A = A - BR	00010010		
		rightShift	0000 1001	0111 0110	4
0	0	rightShift	0000 0100	1011 1011	3
1	0	A = A + BR	10110111		
		rightShift	1101 1011	1101 1101	2
1	0	A = A - BR	00001110		
		rightShift	0000 0111	0110 1110	1
0	0	rightShift	0000 0011	1011 0111	0

decimal value. Then the two values is added together.
It can be seen that overflow does occur.

Consider the following two single-precision floating point numbers following the IEEE 754 standard:

C=0 1000 0011 1011 0101 0000000000000000

D=0 1000 1001 0011 1010 0000000000000000

1.1.5 What is the decimal value of C and D ?

C=1.

1.2 Question

In the following questions, when writing the negation of a Boolean expression you may speedup writing by using, e.g., A' instead of \overline{A} , $(AB)'$ instead of \overline{AB} , and $A'B'$ instead of \overline{AB} . Though, be careful with correct grouping of expressions.

Consider the following Karnaugh map for a Boolean function F . For some inputs there is no requirement for a specific output, indicated by d ("don't care")

		CD			
		00	01	11	10
AB	00	d			d
	01		1	d	
	11		1	1	
	10	1	d	1	

1.2.1 Write the truth table for F

Using the Karnaugh map a truth table with the 16 possible scenarios can be written out from the four variables A,B,C, and D

1.2.2 Based on the Karnaugh map, write a minimal Boolean expression for F .

The karnaugh maps can be circled by two circles four sized circles around the two middle groups of ones and d's.

Lastly a circle of size two can be done around the one in the down right corner and the d next to it.

This will create the expression

$$F = AD + BD + A(BC)'$$

A	B	C	D	F
0	0	0	0	d
1	0	0	0	1
0	1	0	0	0
1	1	0	0	0
0	0	1	0	d
1	0	1	0	0
0	1	1	0	0
1	1	1	0	0
0	0	0	1	0
1	0	0	1	d
0	1	0	1	1
1	1	0	1	1
0	0	1	1	0
1	0	1	1	1
0	1	1	1	d
1	1	1	1	1

1.3 Question

A byte-addressable machine with 24-bit addresses has a 2-way set associative cache, where each line can hold a block of 16 bytes. At this time the cache is in the following state, where the tags are written in hexadecimal.

Set	Way 1			Way 2		
	Valid	Tag	Data	Valid	Tag	Data
0	1	6fae	⟨...⟩	1	729e	⟨...⟩
1	1	46fb	⟨...⟩	1	69bb	⟨...⟩
2	1	8ac1	⟨...⟩	0	9c12	⟨...⟩
3	0	cf04	⟨...⟩	0	03f9	⟨...⟩
4	1	1b4c	⟨...⟩	1	8b83	⟨...⟩
5	0	5d96	⟨...⟩	0	a9e8	⟨...⟩
6	1	13af	⟨...⟩	0	6c4e	⟨...⟩
7	0	aa95	⟨...⟩	1	a928	⟨...⟩
8	0	dc17	⟨...⟩	1	cc72	⟨...⟩
9	0	419a	⟨...⟩	0	a99b	⟨...⟩
10	1	fbff	⟨...⟩	1	568d	⟨...⟩
11	1	f080	⟨...⟩	0	4bee	⟨...⟩
12	1	4f97	⟨...⟩	1	5208	⟨...⟩
13	1	f21b	⟨...⟩	1	be46	⟨...⟩
14	0	9336	⟨...⟩	1	5f70	⟨...⟩
15	0	1f94	⟨...⟩	0	5f43	⟨...⟩

1.3.1 Describe the address format for using the cache

Offset / Word field = $\log(\text{block size}) = \log(16) = 4 \text{ bit}$

Index / Set field = $\log(\text{set size}) = \log(16) = 4 \text{ bit}$

Line numbers: 24

Tag = Adresse space - Offset - Index = $24 - 4 - 4 = 16$ making the address format:

Tag 16 bit, Set 4 bit, Word 4 bit

1.3.2 Imagine 6 alternate realities where the next memory request is to read or write to each of the following addresses (written here in hexadecimal). For each scenario, describe whether there is a hit or miss in the cache.

1. 568da3 - set: 10, tag: 568d - Hit
2. 419a99 - set: 9, tag: 419a - Not valid
3. 1f005d - set: 5, tag: 1f00 - Miss
4. 69bb1b - set: 1, tag: 69bb - Hit
5. ab61ca - set: 1, tag: ab61 - Miss
6. 1b4c4e - set: 4, tag: 1b4c - Hit

1.3.3 The cache has multiple lines per set, so the depiction above does not show all fields. Assume the replacement algorithm used in the cache is LRU, which maintains an additional bit per cache line. Compared to the actual data stored in the cache, how many percent additional memory is used in the cache for all required meta-data?

For LRU a bit for each line is needed.

Therefore the LRU is $1 \cdot 16 = 16 \text{ bit}$

The cache itself is 16 lines of 4 bit set, 2 times (1 bit valid, 16 bit tag) therefore $16 \cdot (4 + 2(1 + 16)) = 608$

In total the meta data is $16 + 608 = 624 \text{ bit}$.

The cache data is two blocks of 16 bytes pr line $16 \cdot 16 \cdot 2 = 512 \text{ byte}$
 $= 256 \cdot 8 = 4096 \text{ bit}$

Therefore making the meta data being $\frac{624}{4096} = 15.23\%$

- 1.3.4** We would like to use this cache design for the cache in each processor in an symmetric multi-processing system. However, the caches are to be kept coherent using the MESI protocol. Which changes to the cache organisation are needed, if any?

It can be adopted by making the validation 1 bit larger giving it the four possible combinations needed for MESI

1.4 Question

You find yourself debugging a program which sends and receives commands over the Internet. Due to the unreliability of the network connection the program uses error-correction codes for messages. Specifically, each 8-bit message is encoded as a 12-bit Hamming code

- 1.4.1** The program would like to send the following 8 bits of data, written from left to right as d_1, d_2, \dots, d_8 : 1111 0001. Compute the corresponding 12-bit Hamming code, and indicate the check bits.

The check bits is calculated as, where an odd number of ones is 1 and 0 if even

$$C1 = D1, D2, D4, D5, D7 = 1, 1, 1, 0, 0 = 1$$

$$C2 = D1, D3, D4, D6, D7 = 1, 1, 1, 0, 0 = 1$$

$$C4 = D2, D3, D4, D8 = 1, 1, 1, 1 = 0$$

$$C8 = D5, D6, D7, D8 = 0, 0, 0, 1 = 1$$

Making the whole data bit:

$$C1, C2, D1, C4, D2, D3, D4, C8, D5, D6, D7, D8$$

$$1\ 1\ 1\ 0\ 1\ 1\ 1\ 1\ 0\ 0\ 0\ 1$$

- 1.4.2** The program is receives the following 12 bits in response, written from left to right as b_1, b_2, \dots, b_{12} : 1001 1100 0001. Assuming that at most 1 single-bit error has occurred, compute the 8-bit message.

$C1, C2, D1, C4, D2, D3, D4, C8, D5, D6, D7, D8$ 1 0 0 1 1 1 0 0 0 0 0 1 Calculating the check bits

$$C1 = D1, D2, D4, D5, D7 = 0, 1, 0, 0, 0 = 1 \text{ Match}$$

$$C2 = D1, D3, D4, D6, D7 = 0, 1, 0, 0, 0 = 0 \text{ Wrong}$$

$$C4 = D2, D3, D4, D8 = 1, 1, 0, 1 = 1 \text{ Match}$$

$$C8 = D5, D6, D7, D8 = 0, 0, 0, 1 = 0 \text{ Wrong}$$

Since the C2 and C8 share D6 it must be flipped making the original message:
0110 0101

2 2018

2.1 Question

Consider the two 5 bit patterns $s = 10011$ and $T = 01010$.

2.1.1 Interpret S and T as unsigned integers. What are their decimal values?

$$S = 2^5 + 2^1 + 2^0 = 19$$

$$T = 2^4 + 2^1 = 10$$

2.1.2 Interpret S and T as signed integers in two's complement representation. What are their decimal values?

$$S = -2^5 + 2^1 + 2^0 = -13$$

$$T = 2^4 + 2^1 = 10$$

The first bit will then be used as a negative value when calculated, and since the first value is already zero in T the value does not change.

2.1.3 Interpret S and T as signed integers in two's complement representation. Multiply S and T using Booth's algorithm. Show

- The initial state of the algorithm
- For each step the actions taken the resulting state
- The final product in two's complement representation and its decimal value

$$\text{Result} = 1101111110 = -2^9 + 2^8 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 = -130$$

Consider 9-bit signed integers in two's complement representation

multiplicand: 10011 Multiplier: 01010					
Q[0]	Q-1	Log	A	Q	SC
0	0	initial	0 0000	0 1010	5
0	0	rightShift	0 0000	0 0101	4
1	0	A = A - M	0 1101	0 0101	
		rightShift	0 0110	1 0010	3
0	1	A = A + M	1 1001	1 0010	
		rightShift	1 1100	1 1001	2
1	0	A = A - M	0 1001	1 1001	
		rightShift	0 0100	1 1100	1
0	1	A = A + M	1 0111	1 1100	
		rightShift	1 1011	1 1110	0

multiplicand: 01010 Multiplier: 10011					
Q[0]	Q-1	Log	A	Q	SC
1	0	initial	0 0000	1 0011	5
1	0	A = A - M	1 0110	1 0011	
		rightShift	1 1011	0 1001	4
1	1	rightShift	1 1101	1 0100	3
0	1	A = A + M	0 0111	1 0100	
		rightShift	0 0011	1 1010	2
0	0	rightShift	0 0001	1 1101	1
1	0	A = A - M	1 0111	1 1101	
		rightShift	1 1011	1 1110	0

2.1.4 What is the range of integers are representable in this format?

The lowest value is $-2^8 = -256$ where the most significant bit is 1 and the rest is zero. The highest will be everything is one except the most significant bit $2^7 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0 = 255$

2.1.5 Which bit pattern represents the lowest and highest representable values?

Highest: 10000000

Lowest: 01111111

Consider the following two single-precision floating point numbers following the IEEE 754 standard:

A= 0 0111 0001 0011 0100 0000000000000000

$B = 0\ 0111\ 0011\ 1001\ 1000\ 0000000000000000$

2.1.6 Give the decimal value of A and B in the normalised form $S \cdot 2^E$ where $1 \leq S \leq 2$

A

Sign: +

Power: $11100010 - 127 = 99$

Decimal: $1.01101000000000 = 1.40625$

$1.40625 \cdot 2^{99}$

B

Sign: +

Power: $01110011 - 127 = -12$

Decimal: $1100110000000000 = 1.59375$

$1.59375 \cdot 2^{-12}$

2.1.7 Perform the addition of A and B in binary arithmetic

For the calculation the zeroes is described as (number of zeroes)x0 instead of writing them out.

X exponent: 99

Y exponent: -12

X fraction: 101101000000000000000000

Y fraction: 110011000000000000000000

First align mantisse

Add 111 zeroes to Y to align

Y fraction: (111)x0 110011000000000000000000 Adding X fraction and Y fraction

$101101\ (128)\text{x}0\ 110011000000000000000000$

Sum: $101101\ (105)\text{x}0\ 110011000000000000000000$

Convert sum back into IEEE 754 format, by removing the first 1 if the sum and using the largest exponent

$0\ 11100010\ 011010000000000000000000$

2.2 Question

In the following questions, when writing the negation of a Boolean expression you may speedup writing by using, e.g., A' instead of \overline{A} , $(AB)'$ instead of \overline{AB} , and $A'B'$ instead of \overline{AB} . Though, be careful with correct grouping of expressions.

Consider the following truth table for a Boolean function that we would like

to implement in hardware with the smallest circuit we can reasonably find. This means that we should take advantage of the fact that for some inputs there are no requirement for a specific output, indicated by d (“don’t care”).

A	B	C	D	F
0	0	0	0	1
0	0	0	1	0
0	0	1	0	1
0	0	1	1	0
0	1	0	0	1
0	1	0	1	d
0	1	1	0	d
0	1	1	1	0
1	0	0	0	0
1	0	0	1	d
1	0	1	0	1
1	0	1	1	0
1	1	0	0	d
1	1	0	1	1
1	1	1	0	0
1	1	1	1	0

2.2.1 As a first solution, write a Boolean expression for F in sum-of-products form.

$$F = (ABCD)' + C(ABD)' + B(ACD)' + AC(DB)' + ABCD'$$

2.2.2 Draw the Karnaugh map for F

			CD		
		00	01	11	10
	00	1			1
AB	01	1	d		d
	11	d	1		
	10		d		1

2.2.3 Based on the Karnaugh map, write a minimal Boolean expression for F .

$$F = (AD)' + BC' + C(BD)'$$

2.3 Question

We are tasked with designing a new cache for a machine with 32-bit addresses, where the addressable unit is a single 8-bit byte. The size of each cache line must be able to hold 64 bytes of data, and in total the cache must be of size 128 KiB (where 1KiB is 1024B). First we consider implementing the cache with direct mapping

2.3.1 How many lines will the cache have?

Physical address bits = 32 bit

Block size = 64 bytes = 2^6 B

Cache size = $128 \cdot 1024 = 2^{17}$ B

Number of lines = Cache size/Block size = 2^{11}

2.3.2 Describe the address format for the cache?

Block offset = $\log(\text{Block size}) = 6$ bits

Line offset = $\log(\text{Number of lines}) = 11$ bits

Tag size = $32 - 6 - 11 = 15$ bits

Making the address format: Tag 15 bits — Line offset 11 bits — Block offset 6 bits

2.3.3 Besides the actual data, for each line we must also store a tag and a valid-bit. Compared to the memory used for data, how many percent additional memory do we need for the required meta-data?

The tag was found to be 15 bits and a valid bit making the meta data 16 bit pr. line.

This making the meta data take up $\frac{16b}{64.8b} = 3.12\%$ of the stored data.

Alternatively we now consider an implementation as a 8-way set associative cache

2.3.4 How many sets will the cache have?

The number of sets will be = number of lines/8=2³

2.3.5 Describe the address format for the cache?

The block offset is still 6 bits

The set offset will be = log(number of sets)=3 bits

Making the tag = 32-6-3=23 bits

2.3.6 Besides the actual data, for each line we must also store a tag and a valid-bit. Compared to the memory used for data, how many percent additional memory do we need for the required meta-data?

The tag is now 23 bits so the meta data is 24 bits.

This making the meta data take up $\frac{24b}{64.8b} = 4.69\%$ of the stored data.

The two cache designs are now compared on the effect on memory access times.

2.3.7 The direct-mapped cache has an access time for each memory request of 4ns and a miss rate of 5%. After a miss the request goes to main memory, which adds 103ns to the access time (including the subsequent cache update). What is the average access time?

$$4ns \cdot 0.95 + 103ns \cdot 0.05 = 8.95ns$$

2.3.8 Due to its increased complexity the set-associative cache has a slightly higher accesstime of 6ns, though the miss rate drops to just 3%. As the cache update in case of a miss is also slightly slower, the added access time is now 104ns. What is the averageaccess time for this setup?

$$6ns \cdot 0.97 + 104ns \cdot 0.03 = 8.94$$

2.4 Question

Consider a symmetric multi-processing system with two processors P_1 and P_2 . They each have their own cache which are kept coherent using the MESI protocol. For this question we look at two blocks of memory b_1 and b_2 where we assume they are mapped to the caches such that they can be present simultaneously. Initially (at time t_0), the state of b_1 is “invalid” in the cache of both P_1 and P_2 , while the state of b_2 is “exclusive” in P_1 and “invalid” in P_2 .

2.4.1 For the following sequence of events, show the resulting MESI state of each block ineach processor.

Time	Event	P_1		P_2	
		b_1	b_2	b_1	b_2
t_0	(Initial state)	I	E	I	I
t_1	P_1 : read b_1				
t_2	P_2 : read b_1				
t_3	P_2 : read b_2				
t_4	P_1 : read b_2				
t_5	P_1 : read b_1				

t_1 - Read b_1 from memory and sets it as exclusive

t_2 - Read b_1 from memory and set as shared and notifies p_1 to set b_1 as shared

t_3 - Read b_2 from memory and set as shared and notifies p_1 to set b_2 as shared

t_4 - Nothing changes

t_5 - Nothing changes

2.5 Question

Assume the following content in the main memory(all numbers are in hexadecimal format with omitted leading zeros):

Name	Address	Value
	...	
	0x80	0x88
	0x88	0x90
	0x90	0x98
	...	
	0xC0	0xFFF8
	0xC8	0x88
	...	
a:	0xFFF8	0x80
	...	

During the course of your program, you have defined the local variables a and b on the stack as follows:

```
01 | ...
02 | long **a
03 | ...
```

Assume you are working on a 64 Bit Linux machine using the gcc;thus long values as well as addresses consume 64 bit.The execution of your program has led to the memory content as displayed above.

2.5.1 To what values do the following expressions evaluate

- `**a - 0x80`
- `&a - 0xFFF8`
- `**(a+1) - 0xFFF0`

2.5.2 What is a void pointer and what is such a pointer used for?

A void pointer is a pointer which points to any data type.

A void pointer can be used as a generic pointer.

2.5.3 Declare a function pointer “ptr” to a function requiring two integers as arguments returning a pointer to a long variable.

```
01 | long func(int a, int b);  
02 |  
03 | long (*ptr)(int, int) = &func
```

Now consider the following C function calculating the scalar product of two vectors a and b of dimension d. The program is written for a 32 Bit Linux machine.

```
01 | int scalarproduct(int* a, int* b, int d)  
02 | {  
03 |     int ret = 0;  
04 |     int i;  
05 |     for(i = 0; i < d; i++){  
06 |         ret += a[i]*b[i];  
07 |     }  
08 |     return ret;  
09 | }
```

2.5.4 Give the stack-frame relevant for the function relative to the function’s stack base pointer (%ebp). Identify all variables and parameters of the C-Function. Use the following table as template (produce a table similar to this one):

Location in stack	Content
...	
-8(%ebp)	Local variable
-4(%ebp)	Local variable
(%ebp)	Old base pointer
4(%ebp)	Return address
8(%ebp)	function arguments
...	

```

1      .type    scalarproduct, @function
2      scalarproduct:
3          pushl %ebp
4          movl  %esp, %ebp
5          subl  $16, %esp
6          movl  $0, -8(%ebp)
7          movl  $0, -4(%ebp)
8          jmp   .L2
9      .L3:
10         movl  -4(%ebp), %eax
11         leal  0(,%eax,4), %edx
12         movl  8(%ebp), %eax
13         addl  %edx, %eax
14         movl  (%eax), %edx
15         movl  -4(%ebp), %eax
16         leal  0(,%eax,4), %ecx
17         movl  12(%ebp), %eax
18         addl  %ecx, %eax
19         movl  (%eax), %eax
20         imull %edx, %eax
21         addl  %eax, -8(%ebp)
22         addl  $1, -4(%ebp)
23     .L2:
24         movl  -4(%ebp), %eax
25         cmpl  16(%ebp), %eax
26         jl    .L3
27         movl  -8(%ebp), %eax
28         movl  %ebp,%esp
29         pop   %ebp
30         ret

```

2.5.5 Identify the lines in the assembly code where the abort condition $i < d$ is checked and where does the code continue in case the condition does not hold and where in case it does not.

The comparison is done at line 25, and if less jump to l3 otherwise it continues to line 27.

2.5.6 Identify in which lines in the assembler code the address of $a[i]$ is calculated and in which register is the value of $a[i]$ temporarily stored.

$a[i]$ is calculated at line 13 where first the address of a is found and then the value of i is added together.

The result is then stored in `%edx`

2.5.7 Consider line 16 in the assembly code

- What does this line do? - The line takes the address of `eax` multiplies it by 4 and puts it into `%ecx`
- Replace this line by standard arithmetic or logical operation - `imull 4, (%eax); movl(%eax), %ecx`

2.5.8 Through which register is the result returned to the caller?

Throughout the program the value is stored in `-8(%ebp)` and then moved into `%eax` at line 27 and returned